

Database Systems

SQL

Based on slides by Feifei Li, University of Utah

The SQL Query Language

- SQL stands for Structured Query Language
- The most widely used relational query language. Current standard is SQL:2016
 - (actually there is a new standard with small modifications that has been release in 2019)
 - Many systems like MySQL/PostgreSQL have some “unique” aspects
 - as do most systems.
- Here we concentrate on SQL-92 and SQL:1999

DDL – Create Table

- CREATE TABLE *table_name* ({ *column_name data_type* [DEFAULT *default_expr*] [*column_constraint*[, ...]] | *table_constraint* } [, ...])

- Data Types include:
 - character(n) – fixed-length character string (CHAR(n))
 - character varying(n) – variable-length character string (VARCHAR(n))
 - smallint, integer, bigint, numeric, real, double precision
 - date, time, timestamp, ...
 - serial - unique ID for indexing and cross reference
 - ...
 - you can also define your own type!! (SQL:1999)

Create Table (w/column constraints)

- CREATE TABLE *table_name* ({ *column_name data_type* [DEFAULT *default_expr*] [*column_constraint*[, ...]] | *table_constraint* } [, ...])

Column Constraints:

- [CONSTRAINT *constraint_name*] { NOT NULL | NULL | UNIQUE | PRIMARY KEY | CHECK (*expression*) | REFERENCES *reftable* (*refcolumn*) [ON DELETE *action*] [ON UPDATE *action*] }

action is one of:

NO ACTION, CASCADE, SET NULL, SET DEFAULT

expression for column constraint must produce a boolean result and reference the related column's value only.

Create Table (w/table constraints)

- CREATE TABLE *table_name* ({ *column_name data_type* [DEFAULT *default_expr*] [*column_constraint*[, ...]] | *table_constraint*} [, ...])

Table Constraints:

- [CONSTRAINT *constraint_name*]
{ UNIQUE (*column_name* [, ...]) |
PRIMARY KEY (*column_name* [, ...]) |
CHECK (*expression*) |
FOREIGN KEY (*column_name* [, ...]) REFERENCES *reftable* [(*refcolumn* [, ...])] [ON
DELETE *action*] [ON UPDATE *action*] }

Here, *expressions*, *keys*, etc can include multiple columns

Create Table (Examples)

```
CREATE TABLE films (
    code      CHAR(5) PRIMARY KEY,
    title     VARCHAR(40),
    did       DECIMAL(3),
    date_prod DATE,
    kind      VARCHAR(10),
    CONSTRAINT production UNIQUE(date_prod)
    FOREIGN KEY did REFERENCES distributors ON DELETE NO ACTION );
```

```
CREATE TABLE distributors (
    did   DECIMAL(3) PRIMARY KEY,
    name  VARCHAR(40)
    CONSTRAINT con1 CHECK (did > 100 AND name <> ' ') );
```

The SQL DML

- Single-table queries are straightforward.
- To find all 18 year old students, we can write:

```
SELECT *
FROM Students S
WHERE S.age=18
```

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@ee	18	3.2

- To find just names and logins, replace the first line:

```
SELECT S.name, S.login
```

Querying Multiple Relations

- Can specify a join over two tables as follows:

```
SELECT S.name, E.cid  
FROM Students S, Enrolled E  
WHERE S.sid=E.sid AND E.grade='B'
```

sid	cid	grade
53831	Carnatic101	C
53831	Reggae203	B
53650	Topology112	A
53666	History105	B

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@ee	18	3.2

Note: obviously no referential integrity constraints have been used here.

result =

S.name	E.cid
Jones	History105

Basic SQL Query

SELECT	[DISTINCT] target-list
FROM	relation-list
WHERE	qualification

- relation-list : A list of relation names
 - possibly with a *range-variable* after each name
- target-list : A list of attributes of tables in *relation-list*
- qualification : Comparisons combined using AND, OR and NOT.
 - Comparisons are Attr *op* const or Attr1 *op* Attr2, where *op* is one of
 $<$, $>$, $=$, \leq , \geq , \neq
- DISTINCT: optional keyword indicating that the answer should not contain duplicates.
 - In SQL SELECT, the default is that duplicates are not eliminated! (Result is called a “multiset”)

Query Semantics

- Semantics of an SQL query are defined in terms of the following conceptual evaluation strategy:
 1. do FROM clause: compute *cross-product* of tables (e.g., Students and Enrolled).
 2. do WHERE clause: Check conditions, discard tuples that fail. (called “*selection*”).
 3. do SELECT clause: Delete unwanted fields. (called “*projection*”).
 4. If DISTINCT specified, eliminate duplicate rows.
- Probably the least efficient way to compute a query!
 - An optimizer will find more efficient strategies to get the *same answer*.

Step 1 – Cross Product

S.sid	S.name	S.login	S.age	S.gpa	E.sid	E.cid	E.grade
53666	Jones	jones@cs	18	3.4	53831	Carnatic101	C
53666	Jones	jones@cs	18	3.4	53832	Reggae203	B
53666	Jones	jones@cs	18	3.4	53650	Topology112	A
53666	Jones	jones@cs	18	3.4	53666	History105	B
53688	Smith	smith@ee	18	3.2	53831	Carnatic101	C
53688	Smith	smith@ee	18	3.2	53831	Reggae203	B
53688	Smith	smith@ee	18	3.2	53650	Topology112	A
53688	Smith	smith@ee	18	3.2	53666	History105	B

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@ee	18	3.2

X

sid	cid	grade
53831	Carnatic101	C
53831	Reggae203	B
53650	Topology112	A
53666	History105	B

SELECT S.name, E.cid
 FROM Students S, Enrolled E
 WHERE S.sid=E.sid AND E.grade='B'

Step 2 - Discard tuples that fail predicate

S.sid	S.name	S.login	S.age	S.gpa	E.sid	E.cid	E.grade
53666	Jones	jones@cs	18	3.4	53831	Carnatic101	C
53666	Jones	jones@cs	18	3.4	53832	Reggae203	B
53666	Jones	jones@cs	18	3.4	53650	Topology112	A
53666	Jones	jones@cs	18	3.4	53666	History105	B
53688	Smith	smith@ee	18	3.2	53831	Carnatic101	C
53688	Smith	smith@ee	18	3.2	53831	Reggae203	B
53688	Smith	smith@ee	18	3.2	53650	Topology112	A
53688	Smith	smith@ee	18	3.2	53666	History105	B

```
SELECT S.name, E.cid  
FROM Students S, Enrolled E  
WHERE S.sid=E.sid AND E.grade='B'
```

Step 3 - Discard Unwanted Columns

S.sid	S.name	S.login	S.age	S.gpa	E.sid	E.cid	E.grade
53666	Jones	jones@cs	18	3.4	53831	Carnatic101	C
53666	Jones	jones@cs	18	3.4	53832	Reggae203	B
53666	Jones	jones@cs	18	3.4	53650	Topology112	A
53666	Jones	jones@cs	18	3.4	53666	History105	B
53688	Smith	smith@ee	18	3.2	53831	Carnatic101	C
53688	Smith	smith@ee	18	3.2	53831	Reggae203	B
53688	Smith	smith@ee	18	3.2	53650	Topology112	A
53688	Smith	smith@ee	18	3.2	53666	History105	B

```
SELECT S.name, E.cid  
FROM Students S, Enrolled E  
WHERE S.sid=E.sid AND E.grade='B'
```

Now the Details

We will use these instances of relations in our examples.

Question:

If the key for the Reserves relation contained only the attributes *sid* and *bid*, how would the semantics differ?

Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

Sailors

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

Boats

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

Example Schemas

```
CREATE TABLE Sailors (sid INTEGER PRIMARY KEY,  
    sname CHAR(20), rating INTEGER, age REAL)
```

```
CREATE TABLE Boats (bid INTEGER PRIMARY KEY,  
    bname CHAR (20), color CHAR(10))
```

```
CREATE TABLE Reserves (  
    sid INTEGER REFERENCES Sailors,  
    bid INTEGER, day DATE,  
    PRIMARY KEY (sid, bid, day),  
    FOREIGN KEY (bid) REFERENCES Boats)
```

Another Join Query

```
SELECT sname  
FROM Sailors, Reserves  
WHERE Sailors.sid=Reserves.sid  
      AND bid=103
```

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
95	Bob	3	63.5	22	101	10/10/96
95	Bob	3	63.5	95	103	11/12/96

Some Notes on Range Variables

- Can associate “range variables” with the tables in the FROM clause.
 - saves writing, makes queries easier to understand
- Needed when ambiguity could arise.
 - for example, if same table used multiple times in same FROM (called a “self-join”)

```
SELECT sname  
FROM Sailors,Reserves  
WHERE Sailors.sid=Reserves.sid AND bid=103
```

Can be rewritten using range variables as:

```
SELECT S.sname  
FROM Sailors S, Reserves R  
WHERE S.sid=R.sid AND bid=103
```

More Notes

- Here's an example where range variables are required (self-join example):

```
SELECT x.sname, x.age, y.sname, y.age  
FROM Sailors x, Sailors y  
WHERE x.age > y.age
```

- Note that target list can be replaced by “*” if you don't want to do a projection:

```
SELECT *  
FROM Sailors x  
WHERE x.age > 20
```

Find sailors who've reserved at least one boat

```
SELECT S.sid  
      FROM Sailors S, Reserves R  
     WHERE S.sid=R.sid
```

- Would adding DISTINCT to this query make a difference (DISTINCT forces the system to remove duplicates from the output)?
- What is the effect of replacing *S.sid* by *S.sname* in the SELECT clause?
 - Would adding DISTINCT to this variant of the query make a difference?

Expressions

- Can use arithmetic expressions in SELECT clause (plus other operations we'll discuss later)
- Use **AS** to provide column names (like a renaming operator)

```
SELECT S.age, S.age-5 AS age1, 2*S.age AS age2  
      FROM Sailors S  
 WHERE S.sname = 'Dustin'
```

- Can also have expressions in WHERE clause:

```
SELECT S1.sname AS name1, S2.sname AS name2  
      FROM Sailors S1, Sailors S2  
 WHERE 2*S1.rating = S2.rating - 1
```

String operations

- SQL supports some basic string operations: “LIKE” is used for string matching

```
SELECT S.age, S.age-5 AS age1, 2*S.age AS age2  
      FROM Sailors S  
 WHERE S.sname LIKE 'J_%m'
```

'_' stands for any one character and '%' stands for 0 or more arbitrary characters.

Find sid's of sailors who've reserved a red or a green boat

- UNION: Can be used to compute the union of any two *union-compatible* sets of tuples (which are themselves the result of SQL queries).

```
SELECT R.sid  
      FROM Boats B,Reserves R  
     WHERE R.bid=B.bid AND  
(B.color='red' OR B.color='green')
```

vs.

```
SELECT R.sid  
      FROM Boats B, Reserves R  
     WHERE R.bid=B.bid AND B.color='red'  
UNION  
SELECT R.sid  
      FROM Boats B, Reserves R  
     WHERE R.bid=B.bid AND B.color='green'
```

Find sid's of sailors who've reserved a red and a green boat

- If we simply replace OR by AND in the previous query, we get the wrong answer. (Why?)
- Instead, could use a self-join:

```
SELECT R1.sid  
FROM Boats B1, Reserves R1,  
      Boats B2, Reserves R2  
WHERE R1.sid=R2.sid  
      AND R1.bid=B1.bid  
      AND R2.bid=B2.bid  
      AND (B1.color='red' AND B2.color='green')
```

Find sid's of sailors who've reserved a red and a green boat

- Or you can use **AS** to “rename” the output of a SQL block:

```
SELECT R1.sid  
FROM Boats B1, Reserves R1,  
      (SELECT R2.sid  
       FROM Boats B2, Reserves R2  
      WHERE B2.color ='green'  
        AND B2.bid=R2.bid) AS GR  
WHERE R1.sid=GR.sid  
  AND R1.bid=B1.bid  
  AND B1.color='red'
```

```
SELECT RR.sid  
FROM (SELECT R1.sid  
      FROM Boats B1, Reserves R1,  
      WHERE B1.color='red'  
        AND B1.bid=R1.bid) AS RR,  
      (SELECT R2.sid  
       FROM Boats B2, Reserves R2  
      WHERE B2.color ='green'  
        AND B2.bid=R2.bid) AS GR  
WHERE RR.sid=GR.sid
```

AND Continued...

- **INTERSECT**: Can be used to compute the intersection of any two *union-compatible* sets of tuples.
- **EXCEPT** (sometimes called MINUS)
- **many** systems don't support them.

```
SELECT S.sid ← Key field!
FROM Sailors S, Boats B,
      Reserves R
WHERE S.sid=R.sid
      AND R.bid=B.bid
      AND B.color='red'
INTERSECT
SELECT S.sid
FROM Sailors S, Boats B,
      Reserves R
WHERE S.sid=R.sid
      AND R.bid=B.bid
      AND B.color='green'
```

Find sid's of sailors who've reserved a red but did not reserve a green boat

```
SELECT S.sid  
FROM Sailors S, Boats B,  
      Reserves R  
WHERE S.sid=R.sid  
      AND R.bid=B.bid  
      AND B.color='red'  
EXCEPT  
SELECT S.sid  
FROM Sailors S, Boats B,  
      Reserves R  
WHERE S.sid=R.sid  
      AND R.bid=B.bid  
      AND B.color='green'
```

Nested Queries

- Powerful feature of SQL: WHERE clause can itself contain an SQL query!
 - Actually, so can FROM and HAVING clauses.

Names of sailors who've reserved boat #103:

```
SELECT S.sname  
FROM Sailors S  
WHERE S.sid IN (SELECT R.sid  
                 FROM Reserves R WHERE R.bid=103)
```

- To find sailors who've *not* reserved #103, use NOT IN.
- To understand semantics of nested queries:
 - think of a *nested loops* evaluation: *For each Sailors tuple, check the qualification by computing the subquery.*

Nested Queries with Correlation

Find names of sailors who've reserved boat #103:

```
SELECT S.sname  
FROM Sailors S  
WHERE EXISTS (SELECT *  
              FROM Reserves R  
              WHERE R.bid=103 AND S.sid=R.sid)
```

- **EXISTS** is another set comparison operator, like **IN**.
- Can also specify **NOT EXISTS**
- If **UNIQUE** is used, and * is replaced by *R.bid*, finds sailors with at most one reservation for boat #103.
 - **UNIQUE** checks for duplicate tuples in a subquery;
 - **UNIQUE** returns true for empty subquery (assumes that two NULL values are different)
- Subquery must be recomputed for each *Sailors* tuple.
 - Think of subquery as a function call that runs a query!

More on Set-Comparison Operators

- We've already seen IN, EXISTS and UNIQUE. Can also use NOT IN, NOT EXISTS and NOT UNIQUE.
- Also available: *op ANY*, *op ALL*
- Find sailors whose rating is greater than that of some sailor called Horatio:

```
SELECT *
FROM Sailors S
WHERE S.rating > ANY (SELECT S2.rating FROM Sailors S2
                      WHERE S2.sname='Horatio')
```

Semantics of nested operators

- v is a value, A is a *multi-set*
- $v \text{ IN } A$ evaluates to true iff $v \in A$. $v \text{ NOT IN } A$ is the opposite.
- $\text{EXISTS } A$ evaluates to true iff $A \neq \emptyset$. $\text{NOT EXISTS } A$ is the opposite.
- $\text{UNIQUE } A$ evaluates to true iff A is a *set*. $\text{NOT UNIQUE } A$ is the opposite.
- $v \text{ OP ANY } A$ evaluates to true iff $\exists x \in A$, such that $v \text{ OP } x$ evaluates to true.
- $v \text{ OP ALL } A$ evaluates to true iff $\forall x \in A$, $v \text{ OP } x$ *always* evaluates to true.

Rewriting INTERSECT Queries Using IN

Find sid's of sailors who've reserved both a red and a green boat:

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid
    AND B.color='red'
    AND R.sid IN (SELECT R2.sid
                  FROM Boats B2, Reserves R2
                  WHERE R2.bid=B2.bid
                      AND B2.color='green')
```

- Similarly, EXCEPT queries re-written using NOT IN.
- How would you change this to find *names* (not *sid's*) of Sailors who've reserved both red and green boats?

Division in SQL (For All query)

Find sailors who've reserved all boats.

```
SELECT S.sname  
FROM Sailors S      Sailors S such that ...  
WHERE NOT EXISTS (SELECT B.bid   there is no boat B  
                   FROM Boats B   without ...  
                   WHERE NOT EXISTS (SELECT R.bid  
                                     FROM Reserves R  
                                     WHERE R.bid=B.bid  
                                           AND R.sid=S.sid))  
  
a Reserves tuple showing S reserved B
```

Division in SQL (For All query) Another way..

Find sailors who've reserved all boats.

```
SELECT S.sname      Sailors S such that ...
FROM Sailors S
WHERE NOT EXISTS ( (SELECT B.bid      there is no boat B
                     FROM Boats B)      without ...
                  EXCEPT
                  (SELECT R.bid      a Reserves
                     FROM Reserves R
                     WHERE R.sid=S.sid))  tuple showing
                                         S reserved B
```

Basic SQL Queries - Summary

- An advantage of the relational model is its well-defined query semantics.
- SQL provides functionality close to that of the basic relational model.
 - some differences in duplicate handling, null values, set operators, etc.
- Typically, many ways to write a query
 - the system is responsible for figuring a fast way to actually execute a query regardless of how it is written.
- Lots more functionality beyond these basic features

Aggregate Operators

- Significant extension from set based queries.

```
SELECT COUNT (*)  
FROM Sailors S
```

```
SELECT AVG (S.age)  
FROM Sailors S  
WHERE S.rating=10
```

```
SELECT COUNT (DISTINCT S.rating)  
FROM Sailors S  
WHERE S.sname='Bob'
```

```
COUNT (*)  
COUNT ( [DISTINCT] A)  
SUM ( [DISTINCT] A)  
AVG ( [DISTINCT] A)  
MAX (A)  
MIN (A)
```

single column

```
SELECT AVG ( DISTINCT S.age)  
FROM Sailors S  
WHERE S.rating=10
```

Find name and age of the oldest sailor(s)

- The first query is incorrect!
- Third query equivalent to second query.

```
SELECT S.sname, MAX (S.age)  
FROM Sailors S
```

```
SELECT S.sname, S.age  
FROM Sailors S  
WHERE S.age =  
(SELECT MAX (S2.age)  
FROM Sailors S2)
```

```
SELECT S.sname, S.age  
FROM Sailors S  
WHERE S.age >= ALL (SELECT S2.age  
FROM Sailors S2)
```

GROUP BY and HAVING

- So far, we've applied aggregate operators to all (qualifying) tuples.
 - Sometimes, we want to apply them to each of several *groups* of tuples.
- Consider: *Find the age of the youngest sailor for each rating level.*
 - In general, we don't know how many rating levels exist, and what the rating values for these levels are!
 - Suppose we know that rating values go from 1 to 10; we can write 10 queries that look like this (!):

For $i = 1, 2, \dots, 10$:

```
SELECT MIN (S.age)
FROM Sailors S
WHERE S.rating = i
```

Queries With GROUP BY

- To generate values for a column based on groups of rows, use **aggregate** functions in SELECT statements with the GROUP BY clause

```
SELECT      [DISTINCT] target-list
FROM        relation-list
[WHERE      qualification]
GROUP BY  grouping-list
```

The *target-list* contains

- (i) list of column names &
- (ii) terms with aggregate operations (e.g., MIN (*S.age*)).
 - column name list (i) can **contain only** attributes from the *grouping-list*, since the output for each group must represent a consistent value from that group.

Group By Examples

For each rating, find the average age of the sailors

```
SELECT S.rating, AVG (S.age)  
FROM Sailors S  
GROUP BY S.rating
```

For each rating find the age of the youngest sailor with age ≥ 18

```
SELECT S.rating, MIN (S.age)  
FROM Sailors S  
WHERE S.age >= 18  
GROUP BY S.rating
```

Conceptual Evaluation

- The cross-product of *relation-list* is computed first, tuples that fail *qualification* are discarded, ‘unnecessary’ fields are deleted, and the remaining tuples are partitioned into groups by the value of attributes in *grouping-list*.
- One answer tuple is generated per qualifying group.

An illustration

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
```

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.0
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

1. Form cross product

rating	age
1	33.0
7	45.0
7	35.0
8	55.0
10	35.0

2. Delete unneeded rows,
columns; form groups

Answer Table

3. Perform
Aggregation

rating	age
1	33.0
7	35.0
8	55.0
10	35.0

Find the number of reservations for each red boat.

```
SELECT B.bid, COUNT(*) AS numres  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid  
    AND B.color='red'  
GROUP BY B.bid
```

- Grouping over a join of two relations.

An illustration

```
SELECT B.bid, COUNT (*) AS scount  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND B.color='red'  
GROUP BY B.bid
```

b.bid	b.color	r.bid
101	blue	101
102	red	101
103	green	101
104	red	101
101	blue	102
102	red	102
103	green	102
104	red	102

1

b.bid	b.color	r.bid
102	red	102

2

b.bid	scount
102	1

answer

Queries With GROUP BY and HAVING

```
SELECT      [DISTINCT] target-list
FROM        relation-list
WHERE       qualification
GROUP BY    grouping-list
HAVING     group-qualification
```

- Use the HAVING clause with the GROUP BY clause to restrict which group-rows are returned in the result set

Conceptual Evaluation

- Form groups as before.
- The *group-qualification* is then applied to eliminate some groups.
 - Expressions in *group-qualification* must have a *single value per group!*
 - That is, attributes in *group-qualification* must be arguments of an aggregate op **or must also appear** in the *grouping-list*.
- One answer tuple is generated per qualifying group.

Find the age of the youngest sailor with age ≥ 18 , for each rating with at least 2 such sailors

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING COUNT (*) > 1
```

rating	age
1	33.0
7	45.0
7	35.0
8	55.5
10	35.0

2

rating	m-age	count
1	33.0	1
7	35.0	2
8	55.0	1
10	35.0	1

3

sid	sname	rating	age
22	Dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

rating	
7	35.0

Answer relation

Find sailors who've reserved all boats.

```
SELECT S.sname  
FROM Sailors S    Sailors S such that ...  
WHERE NOT EXISTS (SELECT B.bid there is no boat B  
                   FROM Boats B without ...  
                   WHERE NOT EXISTS (SELECT R.bid  
                                     a Reserves tuple showing S reserved B FROM Reserves R  
                                     WHERE R.bid=B.bid  
                                     AND R.sid=S.sid))
```

Find sailors who've reserved all boats.

- Can you do this using Group By and Having?

```
SELECT S.sname  
FROM Sailors S, reserves R  
WHERE S.sid = R.sid  
GROUP BY S.sname, S.sid  
HAVING  
    COUNT(DISTINCT R.bid) =  
        ( Select COUNT (*) FROM Boats)
```

Note: must have both sid and name in the GROUP BY clause. Why?

Sailors

sid	sname	rating	age
1	Frodo	7	22
2	Bilbo	2	39
3	Sam	8	27

```

SELECT S.name
FROM Sailors S, reserves R
WHERE S.sid = R.sid
GROUP BY S.name, S.sid
HAVING COUNT(DISTINCT R.bid) =
      ( Select COUNT (*) FROM Boats)
  
```

sname	sid	bid
Frodo	1	102
Bilbo	2	101
Bilbo	2	102
Frodo	1	102
Bilbo	2	103

sname	sid	bid
Frodo	1	102,102
Bilbo	2	101, 102, 103

Boats

bid	bname	color
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

Reserves

count
3

sid	bid	day
1	102	9/12
2	102	9/12
2	101	9/14
1	102	9/10
2	103	9/13

An Illustration

Find the names of the sailors who've reserved most number of boats
for each rating group

```
SELECT S.sname
FROM Sailors S, reserves R
WHERE S.sid = R.sid
GROUP BY S.sname, S.sid
HAVING
    COUNT(R.bid) =
        ( Select MAX(C) FROM
            (SELECT S1.sid, COUNT(*) AS C FROM
                Sailors S1, reserves R1
                WHERE S1.sid = R1.sid AND S1.rating = S.rating
                GROUP BY S1.sid) )
```

Find the names of the sailors who've reserved most number of boats
for each rating group

```
SELECT S.sname
  FROM Sailors S, reserves R
 WHERE S.sid = R.sid
 GROUP BY S.sname, S.sid
 HAVING
   COUNT(R.bid) >= ALL
     (SELECT COUNT(*) FROM
      Sailors S1, reserves R1
     WHERE S1.sid = R1.sid AND S1.rating = S.rating
     GROUP BY S1.sid)
```

INSERT

```
INSERT [INTO] table_name [(column_list)  
VALUES (value_list)
```

```
INSERT [INTO] table_name [(column_list)  
<select statement>]
```

```
INSERT INTO Boats VALUES ( 105, 'Clipper', 'purple')
```

```
INSERT INTO Boats (bid, color) VALUES (99, 'yellow')
```

You can also do a “bulk insert” of values from one table into another:

```
INSERT INTO TEMP(bid)
```

```
SELECT r.bid FROM Reserves R WHERE r.sid = 22;
```

(must be type compatible)

DELETE & UPDATE

```
DELETE [FROM] table_name  
[WHERE qualification]
```

```
DELETE FROM Boats WHERE color = 'red'
```

```
DELETE FROM Boats b  
WHERE b.bid =  
(SELECT r.bid FROM Reserves R WHERE r.sid = 22)
```

Can also modify tuples using UPDATE statement.

```
UPDATE Boats  
SET Color = "green"  
WHERE bid = 103;
```

Null Values

- Field values in a tuple are sometimes *unknown* (e.g., a rating has not been assigned) or *inapplicable* (e.g., no spouse's name).
 - SQL provides a special value *null* for such situations.
- The presence of *null* complicates many issues. E.g.:
 - Special operators needed to check if value is/is not *null*. IS NULL/IS NOT NULL
 - Is *rating>8* true or false when *rating* is equal to *null*? What about AND, OR and NOT connectives?
 - We need a 3-valued logic (true, false and *unknown*).
 - Meaning of constructs must be defined carefully. (e.g., WHERE clause eliminates rows that don't evaluate to true.)
 - New operators (in particular, *outer joins*) possible/needed.

NULLs

e.g. :

branch2=

bname	bcity	assets
Downtown	Boston	9M
Perry	Horse	1.7M
Mianus	Horse	.4M
Kenmore	Boston	NULL

What does this mean?

- We don't know Kenmore's assets?
- Kenmore has no assets?
-

Effect on Queries:

**SELECT * FROM branch2
WHERE assets = NULL**

bname	bcity	assets
Kenmore	Boston	NULL

**SELECT * FROM branch2
WHERE assets IS NULL**

bname	bcity	assets
Kenmore	Boston	NULL

NULLs

- Arithmetic with nulls:

- $n \text{ op null} = \text{null}$
op : +, -, *, /, mod, ...

```
SELECT .....
FROM .....
WHERE boolexpr IS UNKNOWN
```

- Booleans with nulls: One can write:

3-valued logic (true, false, unknown)

What expressions evaluate to UNKNOWN?

1. Comparisons with NULL (e.g. assets = NULL)
2. FALSE OR UNKNOWN (but: TRUE OR UNKNOWN = TRUE)
3. TRUE AND UNKNOWN
4. UNKNOWN AND/OR UNKNOWN

NULLs

Aggregate operations:

```
SELECT SUM(assets)  
FROM branch2
```

returns

SUM

11.1M

NULL is ignored
Same for AVG, MIN, MAX

branch2=

bname	bcity	assets
Downtown	Boston	9M
Perry	Horse	1.7M
Mianus	Horse	.4M
Kenmore	Boston	NULL

But.... COUNT(assets) returns 4!

Let branch3 an empty relation

Then: SELECT SUM(assets)
FROM branch3 returns NULL

but COUNT(<empty rel>) = 0

Joins

```
SELECT (column_list)
FROM table_name
[INNER | {LEFT | RIGHT | FULL } OUTER] JOIN
table_name
ON qualification_list
WHERE ...
```

Explicit join semantics needed unless it is an INNER join (INNER is default)

Inner Join

Only the rows that match the search conditions are returned.

```
SELECT s.sid, S.sname, r.bid  
FROM Sailors s INNER JOIN Reserves r  
ON s.sid = r.sid
```

Returns only those sailors who have reserved boats

```
SELECT s.sid, S.sname, r.bid  
FROM Sailors s NATURAL JOIN Reserves r
```

“NATURAL” means equi-join for each pair of attributes with the same name

An illustration

```
SELECT s.sid, S.sname, r.bid  
FROM Sailors s INNER JOIN Reserves r  
ON s.sid = r.sid
```

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

s.sid	s.name	r.bid
22	Dustin	101
95	Bob	103

Left Outer Join

Left Outer Join returns all matched rows, plus all unmatched rows from the table on the left of the join clause (use nulls in fields of non-matching tuples)

```
SELECT s.sid, S.sname, r.bid  
FROM Sailors s LEFT OUTER JOIN Reserves r  
ON s.sid = r.sid
```

Returns all sailors & information on whether they have reserved boats

An illustration

```
SELECT s.sid, S.sname, r.bid  
FROM Sailors s LEFT OUTER JOIN Reserves r  
ON s.sid = r.sid
```

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

s.sid	s.name	r.bid
22	Dustin	101
95	Bob	103
31	Lubber	

Right Outer Join

Right Outer Join returns all matched rows, plus all unmatched rows from the table on the right of the join clause

```
SELECT r.sid, b.bid, b.name  
FROM Reserves r RIGHT OUTER JOIN Boats b  
ON r.bid = b.bid
```

Returns all boats & information on which ones are reserved.

An illustration

```
SELECT r.sid, b.bid, b.name  
FROM Reserves r RIGHT OUTER JOIN Boats b  
ON r.bid = b.bid
```

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

r.sid	b.bid	b.name
22	101	Interlake
	102	Interlake
95	103	Clipper
	104	Marine

Full Outer Join

Full Outer Join returns all (matched or unmatched) rows from the tables on both sides of the join clause

```
SELECT r.sid, b.bid, b.name  
FROM Reserves r FULL OUTER JOIN Boats b  
ON r.bid = b.bid
```

Returns all boats & all information on reservations

An illustration

```
SELECT r.sid, b.bid, b.name  
FROM Reserves r FULL OUTER JOIN Boats b  
ON r.bid = b.bid
```

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

r.sid	b.bid	b.name
22	101	Interlake
	102	Interlake
95	103	Clipper
	104	Marine

Note: in this case it is the same as the ROJ because bid is a foreign key in reserves, so all reservations must have a corresponding tuple in boats.

Views

```
CREATE VIEW view_name  
AS select_statement
```

Makes development simpler
Often used for security
Not instantiated - makes updates tricky

```
CREATE VIEW Reds  
AS SELECT B.bid, COUNT (*) AS scount  
        FROM Boats B, Reserves R  
       WHERE R.bid=B.bid AND B.color='red'  
      GROUP BY B.bid
```

An illustration

```
CREATE VIEW Reds  
AS SELECT B.bid, COUNT (*) AS scount  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND B.color='red'  
GROUP BY B.bid
```

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

b.bid	scount
102	1

Reds

Views Instead of Relations in Queries

```
CREATE VIEW Reds
AS SELECT B.bid, COUNT (*) AS scount
      FROM Boats B, Reserves R
     WHERE R.bid=B.bid AND B.color='red'
   GROUP BY B.bid
```

bid	scount
102	1

Reds

```
SELECT bname, scount
      FROM Reds R, Boats B
     WHERE R.bid=B.bid
       AND scount < 10
```

Views

Create View vs INTO

(1) `SELECT bname, bcity
 FROM branch
 INTO branch2`

vs

(2) `CREATE VIEW branch2 AS
 SELECT bname, bcity
 FROM branch`

(1) creates new table that gets stored on disk

(2) creates “virtual table” (materialized when needed)

**Therefore: changes in branch are seen in the view version of branch2 (2)
but not for the (1) case.**

Sorting the Results of a Query

- ORDER BY *column* [ASC | DESC] [, ...]

```
SELECT S.rating, S.sname, S.age
      FROM Sailors S, Boats B, Reserves R
     WHERE S.sid=R.sid
           AND R.bid=B.bid AND B.color='red'
 ORDER BY S.rating, S.sname;
```

- Can order by any column in SELECT list, including expressions or aggs, and select top-k:

```
SELECT S.sid, COUNT (*) AS redrescnt
      FROM Sailors S, Boats B, Reserves R
     WHERE S.sid=R.sid
           AND R.bid=B.bid AND B.color='red'
 GROUP BY S.sid
 ORDER BY redrescnt DESC
 LIMIT 10;
```

Discretionary Access Control

GRANT *privileges* ON *object* TO *users* [WITH GRANT OPTION]

- Object can be a **Table** or a **View**
- Privileges can be:
 - Select
 - Insert
 - Delete
 - References (cols) – allow to create a foreign key that references the specified column(s)
 - All
- Can later be **REVOKE**D
- Users can be single users or groups
- See Chapter 17 for more details.

Two more important topics

- Constraints (such as triggers)
- SQL embedded in other languages (not discussed here)
- We will not review them in further details in this class

IC's

What are they?

- predicates on the database
- must always be true (checked whenever db gets updated)

There are the following 4 types of IC's:

Key constraints (1 table)

e.g., 2 accts can't share the same acct_no

Attribute constraints (1 table)

e.g., 2 accts must have nonnegative balance

Referential Integrity constraints (2 tables)

E.g. bnames associated w/ loans must be names of real branches

Global Constraints (n tables)

E.g., a loan must be carried by at least 1 customer with a svngs acct

Global Constraints

Idea: two kinds

- 1) single relation (constraints spans multiple columns)

E.g.: CHECK (total = svngs + check) declared in the CREATE TABLE

- 2) multiple relations: CREATE ASSERTION

SQL examples:

- 1) **single relation: All BOSTON branches must have assets > 5M**

CREATE TABLE branch (

.....

**bcity CHAR(15),
assets INT,
CHECK (NOT(bcity = 'BOS') OR assets > 5M))**

Affects:

insertions into branch

updates of bcity or assets in branch

Global Constraints

SQL example:

- 2) Multiple relations: every loan has a borrower with a savings account

```
CHECK (NOT EXISTS (
    SELECT  *
    FROM    loan AS L
    WHERE   NOT EXISTS(
        SELECT  *
        FROM borrower B, depositor D, account A
        WHERE B cname = D cname AND
              D acct_no = A acct_no AND
              L lno = B lno)))
```

Problem: Where to put this constraint? At depositor? Loan?

Ans: None of the above:

CREATE ASSERTION loan-constraint

CHECK(.....)

**Checked with EVERY DB update!
very expensive.....**

Global Constraints

Issues:

- 1) How does one decide what global constraint to impose?
- 2) How does one minimize the cost of checking the global constraints?

Ans: Semantics of application and Functional dependencies.

Summary: Integrity Constraints

Constraint Type	Where declared	Affects...	Expense
Key Constraints	CREATE TABLE (PRIMARY KEY, UNIQUE)	Insertions, Updates	Moderate
Attribute Constraints	CREATE TABLE CREATE DOMAIN (Not NULL, CHECK)	Insertions, Updates	Cheap
Referential Integrity	Table Tag (FOREIGN KEY REFERENCES)	1. Insertions into referencing rel'n 2. Updates of referencing rel'n of relevant attrs 3. Deletions from referenced rel'n 4. Update of referenced rel'n	1,2: like key constraints. Another reason to index/sort on the primary keys 3,4: depends on <ul style="list-style-type: none"> a. update/delete policy chosen b. existence of indexes on foreign key
Global Constraints	Table Tag (CHECK) or outside table (CREATE ASSERTION)	1. For single rel'n constraint, with insertion, deletion of relevant attrs 2. For assertions w/ every db modification	1. cheap 2. very expensive

Triggers (Active database)

- Trigger: A procedure that starts automatically if specified changes occur to the DBMS
- Analog to a "daemon" that monitors a database for certain events to occur
- Three parts:
 - Event (activates the trigger)
 - Condition (tests whether the triggers should run) [Optional]
 - Action (what happens if the trigger runs)
- Semantics:
 - When event occurs, and condition is satisfied, the action is performed.

An example of Trigger

```
CREATE TRIGGER minSalary BEFORE INSERT ON Professor  
FOR EACH ROW  
WHEN (new.salary < 100,000)  
BEGIN  
    RAISE_APPLICATION_ERROR (-20004, 'Violation of Minimum Professor Salary');  
END;
```

- Conditions can refer to old/new values of tuples modified by the statement activating the trigger.

Triggers – Event,Condition,Action

- Events could be :

BEFORE | AFTER | INSERT | UPDATE | DELETE ON <tableName>

e.g.: BEFORE INSERT ON Professor

- Condition is SQL expression or even an SQL query
(query with non-empty result means TRUE)
- Action can be many different choices :
 - SQL statements , and even DDL and transaction-oriented statements like “commit”.

Example Trigger

Assume our DB has a relation schema :

Professor (pNum, pName, salary)

We want to write a trigger that :

Ensures that any new professor inserted
has salary ≥ 70000

Example Trigger

```
CREATE TRIGGER minSalary BEFORE INSERT ON Professor
```

for what context ?

```
BEGIN
```

check for violation here ?

```
END;
```

Example Trigger

```
CREATE TRIGGER minSalary BEFORE INSERT ON Professor
```

```
FOR EACH ROW
```

```
BEGIN
```

```
Violation of Minimum Professor Salary?
```

```
END;
```

Example Trigger

```
CREATE TRIGGER minSalary BEFORE INSERT ON Professor  
FOR EACH ROW  
  
BEGIN  
  
IF (:new.salary < 70000)  
    THEN RAISE_APPLICATION_ERROR (-20004,  
        'Violation of Minimum Professor Salary');  
END IF;  
  
END;
```

Details of Trigger Example

- BEFORE INSERT ON Professor
 - This trigger is checked before the tuple is inserted
- FOR EACH ROW
 - specifies that trigger is performed for each row inserted
- :new
 - refers to the new tuple inserted
- If (:new.salary < 70000)
 - then an application error is raised and hence the row is not inserted; otherwise the row is inserted.
- Use error code: -20004;
 - this is in the valid range

Example Trigger Using Condition

```
CREATE TRIGGER minSalary BEFORE INSERT ON Professor
FOR EACH ROW
WHEN (new.salary < 70000)
BEGIN
    RAISE_APPLICATION_ERROR (-20004,           'Violation of Minimum
                           Professor Salary');
END;
```

- Conditions can refer to old/new values of tuples modified by the statement activating the trigger.

Triggers: REFERENCING

```
CREATE TRIGGER minSalary BEFORE INSERT ON Professor  
REFERENCING NEW as newTuple  
  
FOR EACH ROW  
  
WHEN (newTuple.salary < 70000)  
  
BEGIN  
    RAISE_APPLICATION_ERROR (-20004,  
        'Violation of Minimum Professor Salary');  
END;
```

Example Trigger

```
CREATE TRIGGER updSalary
    BEFORE UPDATE ON Professor
    REFERENCING OLD AS oldTuple NEW as newTuple
    FOR EACH ROW
    WHEN (newTuple.salary < oldTuple.salary)
    BEGIN
        RAISE_APPLICATION_ERROR (-20004, 'Salary
Decreasing !!!');
    END;
```

- Ensure that salary does not decrease

Another Trigger Example (SQL:99)

```
CREATE TRIGGER youngSailorUpdate
  AFTER INSERT ON SAILORS
  REFERENCING NEW TABLE AS NewSailors
  FOR EACH STATEMENT
  INSERT
    INTO YoungSailors(sid, name, age, rating)
    SELECT sid, name, age, rating
    FROM NewSailors N
    WHERE N.age <= 18
```

Row vs Statement Level Trigger

- Row level: activated once per modified tuple
 - Statement level: activate once per SQL statement
-
- Row level triggers can access new data, statement level triggers cannot always do that (depends on DBMS).
 - Statement level triggers will be more efficient if we do not need to make row-specific decisions

Row vs Statement Level Trigger

- Example: Consider a relation schema

Account (num, amount)

where we will allow creation of new accounts
only during normal business hours.

Example: Statement level trigger

```
CREATE TRIGGER MYTRIG1
BEFORE INSERT ON Account
FOR EACH STATEMENT          --- is default
BEGIN
  IF (TO_CHAR(SYSDATE,'dy') IN ('sat','sun'))
  OR
    (TO_CHAR(SYSDATE,'hh24:mi') NOT BETWEEN '08:00' AND '17:00')
  THEN
    RAISE_APPLICATION_ERROR(-20500,'Cannot      create new account
now !!!');
  END IF;
END;
```

When to use BEFORE/AFTER

- Based on efficiency considerations or semantics.
- Suppose we perform statement-level *after insert*, then all the rows are inserted first, then if the condition fails, and all the inserted rows must be “rolled back”
- Not very efficient !!

Combining multiple events into one trigger

```
CREATE TRIGGER salaryRestrictions
AFTER INSERT OR UPDATE ON Professor
FOR EACH ROW
BEGIN
IF (INSERTING AND :new.salary < 70000) THEN
    RAISE_APPLICATION_ERROR (-20004, 'below min salary'); END IF;
IF (UPDATING AND :new.salary < :old.salary) THEN
    RAISE_APPLICATION_ERROR (-20004, 'Salary Decreasing !!!'); END
    IF;
END;
```

Summary : Trigger Syntax

```
CREATE TRIGGER <triggerName>
BEFORE | AFTER    INSERT | DELETE | UPDATE
[OF <columnList>] ON <tableName> | <viewName>
[REFERENCING [OLD AS <oldName>] [NEW AS <newName>] ]
[FOR EACH ROW] (default is "FOR EACH STATEMENT")
[WHEN (<condition>)]
<PSM body>;
```

Constraints versus Triggers

- Constraints are useful for database consistency
 - Use IC when sufficient
 - More opportunity for optimization
 - Not restricted into insert/delete/update
- Triggers are flexible and powerful
 - Alerters
 - Event logging for auditing
 - Security enforcement
 - Analysis of table accesses (statistics)
 - Workflow and business intelligence ...
- But can be hard to understand
 - Several triggers (Arbitrary order → unpredictable !?)
 - Chain triggers (When to stop ?)
 - Recursive triggers (Termination?)